

External C++ Functionality in CmdStan

Matthew Amato-Yarbrough, mba75@nau.edu

January 2023

Contents

Introduction	2
Installation	2
Repository	2
Bernoulli example	3
Beyond the Bernoulli example	5
Binomial GLM stan model	5
External binomial hpp code	6
External binomial RDump file	7
Final steps	7
Additional points	9
Adding more inputs	9
Including more files	9
Matrices and vertices	9
Limitations	10
Aknowledgements	11

Introduction

The aim of this walk-through is to provide **Stan** users with the knowledge of incorporating external C++ code into their preexisting Stan models. The main benefit of doing so is the ability to create and use custom C++ functions within your Stan model.

The first example of this functionality will use the Bernoulli example code provided by the Stan developers. The code for this example is found in the `/cmdstan/examples/bernoulli` directory of a **CmdStan installation**. The second example will be a binomial model using data from **Linear Models with R**. The code for this section will be provided below.

There exists other Stan interfaces that provide external C++ functionality as well. A notable example is **RStan**. However, the difficulty with using external C++ code is relatively the same in both CmdStan and RStan. The deciding factors in choosing a Stan interface in this case are:

- How comfortable you are with each of the interfaces.
- It is easier to enter data and set hyperparameters within RStan.
- Power users will find CmdStan easier since data entry and hyperparameters are set through the terminal.
- CmdStan is typically more-up-to-date when it comes the Stan back-end.

With any Stan interface, the user will be required to **template** their C++ code to match the templating found in **Stan's math library**. More detail will be provided below regarding the process of templating C++ code to match the templating found in Stan's math library. This requirement remains the biggest time commitment in this process.

Additional information on using external C++ code within RStan can be found in this **vinette**. Using external C++ code within RStan will not be explored in this walk-through.

Installation

The first step is download the version of CmdStan appropriate for your system. The available versions of CmdStan are listed under releases in the **stan-dev/cmdstan** repository. For example: The release that was used for this walk-through is **cmdstan-2.29.2.tar.gz**.

Once the tar.gz is downloaded and unzipped into a directory of your choosing (the system's home directory was used for this example), then you open a terminal window and cd into the unzipped folder (e.g. `cd ~/cmdstan-2.29.2/`). Within your terminal window, run **make build** to complete the CmdStan installation process. If you wanted to delete CmdStan from your system, you can run **make clean-all** from the same directory as you called **make build**.

More recent versions of CmdStan can be used. However, the external C++ functionality may differ with future releases of CmdStan

The details of the system used for this walkthrough:

- Operating System: Ubuntu 20.04 amd64
- CmdStan Version: v2.29.2
- Compiler/Toolkit: g++ 9.4.0 and make 4.2.1

Repository

For detailed terminal output and what you should expect to see after running the **make**, **sample**, and **stansummary** commands, refer to the "terminal_output" directories within the "bernoulli" and "binomial_glm" directories. These directories can be found in the same github repository associated with this documentation: `"/external_cpp_functionality_in_cmdstan"`. This repository also contains all of the code referred to within this walkthrough, as well as examples of CmdStan makefiles that can be used for each set of code.

Bernoulli example

This section will use the Bernoulli source code that is provided by the Stan developers. The associated Stan documentation that expands on the Bernoulli source code using external C++ code can be found [here](#). A condensed version of this documentation will be found below, but it is recommended that you read the Stan documentation if you want additional detail.

First step: Within the Bernoulli directory of your CmdStan installation, open up the file titled “bernoulli.stan”. If you installed CmdStan into your home directory, then this file would be found in `~/cmdstan-2.29.2/examples/bernoulli`.

Above the `data` block, add the following lines of code:

```
functions {  
  real make_odds(real theta);  
}
```

Below the `model` block, add the additional lines of code:

```
generated quantities {  
  real odds;  
  odds = make_odds(theta);  
}
```

This is the essential first step to editing your Stan code so that it can use external C++ code. You are adding a function called `make_odds` that accepts a `real` data type (`theta`) and outputs a `real` data type. Additional information on data types within Stan can be found [here](#). The `make_odds` function will be the function defined within your external C++ file.

Within the same `bernoulli` directory, create a new file titled “`make_odds.hpp`”. In “`make_odds.hpp`”, add the following lines of code:

```
#include <stan/model/model_header.hpp>  
#include <ostream>  
  
namespace bernoulli_model_namespace {  
template <typename T0__,  
          stan::require_all_t<stan::is_stan_scalar<T0__>>* = nullptr>  
          stan::promote_args_t<T0__>  
          make_odds(const T0__& theta, std::ostream* pstream__) {  
  return theta / (1 - theta);  
}  
}
```

For now, save and close “`make_odds.hpp`”. A similar file will be revisited later with additional details.

Open up a terminal window within the CmdStan directory (`~/cmdstan-2.29.2/`) and execute the following line:

```
make STANFLAGS=---allow-undefined USER_HEADER=examples/bernoulli/make_odds.hpp \  
    examples/bernoulli/bernoulli_makeodds
```

In this `make` call, both the `STANFLAGS=---allow-undefined` and `USER_HEADER=` flags are necessary for allowing undefined functions within your Stan code and providing a location to an external header file. As noted by the Stan documentation: “you could put `STANFLAGS` and `USER_HEADER` into the `make/local` file instead of specifying them on the command-line.” This provides users with the option of using a simplified `make` call:

```
make examples/bernoulli/bernoulli_makeodds
```

Assuming that your `make` call was successful, you should find a file titled “`bernoulli_makeodds`” within your `/bernoulli` directory. This is the compiled executable that is a result of your previous `make` call. It will contain the original `bernoulli.stan` model code, as well as the custom external C++ code.

Note: Refer to “`external_cpp_functionality_in_stan/bernoulli/terminal_output/`” on github for an example of a successful `make` call would look like.

Once you have the compiled executable the remaining CmdStan steps are the same.

The Stan developers provide two data files for this example: “`bernoulli.data.json`” and “`bernoulli.data.R`”. Below are the steps one would follow to sampling using the `.json` data file.

Sample using the provided example data: `./bernoulli_makeodds sample data file=bernoulli.data.json`

Display using the `stansummary` functionality: `~/cmdstan-2.29.2/bin/stansummary output.csv`

These steps are discussed in detail in this **part of the CmdStan documentation**. Additional arguments will be used in the next section.

Beyond the Bernoulli example

This section will provide a more complicated example of using external C++ code within CmdStan.

Within the examples directory of your CmdStan installation, you will need to create a new directory as well as several new files. These files will be similar to the files from the previous section. However, the contents of these files will be different. Navigate to `~/cmdstan-2.29.2/examples/` and create a new directory titled “binomial_glm”. Afterwards, open up the newly created directory and create the three following files:

1. binomial_glm.stan
2. external_binomial_glm.hpp
3. bin_glm.Data.list.dump

Note: Depending on your preference, it could be easier to use the **touch** command within terminal to create these files.

Binomial GLM stan model

Open up binomial_glm.stan with a text editor and paste the following contents into the file and then save the file.

```
functions {
  real custom_inv_logit(real value);
}

data {
  int<lower=1> N;           // #obs
  int<lower=1> p;           // #params
  int<lower=1> n[N];        // #trials
  int<lower=0> y[N];        // #successes
  matrix<lower=0, upper=1> [N,p] X; // X matrix
  real bkprior[2];         // bk (k>0) prior mean and sd
}

parameters {
  vector[p] b;
}

model {
  vector[N] output;

  output = X * b;
  for (i in 1:N) {
    // [y | beta]
    y[i] ~ binomial(n[i], custom_inv_logit(output[i]));
  }

  // [beta] = [b0] * prod_k [bk] (k>1)
  b[2:11] ~ normal(bkprior[1], bkprior[2]);

  // [b0 flat improper implied]
}
```

Explanation: The data and parameters block is fairly common for Stan code. The interesting points with the above code are the functions and model blocks.

The functions block: `custom_inv_logit` is the name of the function found in the external C++ file. Similar to the Bernoulli example, the output of the function call is a real number and the input is a real number call value.

The model block: The matrix and vector multiplication is done (`output = X * b`) within the Stan code, rather than within the external C++ code. While this adds additional lines of Stan code, it's to avoid the hassle of matrix and vector multiplication within the external C++ code. Refer to the “Matrices and vertices” subsection below for more information.

External binomial glm code

Now open up `external_binomial_glm.hpp` and paste the following code. Afterwards, save the file.

```
#include "external_binomial_glm.hpp"
#include "stan/math/prim/fun/exp.hpp"

namespace binomial_glm_model_namespace {
template <typename T0__, stan::require_stan_scalar_t<T0__>* = nullptr>
stan::promote_args_t<T0__>
custom_inv_logit(const T0__& value, std::ostream* pstream__)
{
    stan::math::exp_fun exponent_function;

    if (value > 0) {
        return 1.0 / (1.0 + exponent_function.fun(-value));
    }
    else {
        T0__ e = exponent_function.fun(value);
        return e / (1.0 + e);
    }
}
}
```

Explanation: There are a handful of points worth mentioning here.

Inverse logit: This was modeled after Stan's **inv logit struct** found within the Stan library. While this function could be directly used within your Stan model, it was used as the external function within `external_binomial_glm.hpp` for demonstrative purposes.

Exponential function: The exponential function comes from Stan's **exp fun struct**. It's important to note here that you could not get away with using the exponential function from C++'s standard library, as it is not built to accept Stan math types. As such, the best option is to use Stan's `exp_fun` struct. Once the struct option is created (called `exponent_function` above), you can access its exponential function to pass in your `value` parameter.

Creating a new variable: Notice that the variable `e` is of type `T0__`; the same type as the input parameter `value`. This is necessary in order to store the contents of the exponential function.

Initial template: When attempting to compile external C++ code with CmdStan, it will create a separate header file from yours that contains the templated header for your external file. This will happen whether compilation is successful or fails.

- For example: When attempting to run a `make` call with your `binomial_glm.stan` file completed but not your external C++ file, you will encounter a compilation error. But, within the same directory as your `.stan` file, there should be a file called `binomial_glm.hpp`. Open up that file and you should see the appropriate template for your external file towards the top portion of `binomial_glm.hpp`. You will need to include that template (as well as the namespace calls) within your header file.

External binomial RDump file

Now that both the Stan model and external C++ code are ready, all that is left is to run the sampler. However, this will require a data file.

Open up `bin_glm.Data.list.dump`, paste the following lines of code, save and then exit.

```
N <- 24
p <- 11
n <-
c(56, 34, 8, 82, 50, 18, 16, 39, 21, 31, 52, 40, 87, 56,
  20, 35, 23, 9, 45, 72, 49, 28, 36, 24)
y <-
c(9, 19, 7, 5, 6, 12, 5, 21, 19, 7, 14, 25, 58, 45, 20, 11,
  7, 7, 32, 61, 42, 12, 15, 17)
X <-
structure(c(1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
  1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0,
  0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
  1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1,
  0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,
  1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0,
  1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
  0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0,
  0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
  0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
  0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,
  1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
  1, 0, 0, 1),
  .Dim = c(24, 11))
bkprior <-
c(1.77635683940025e-15, 4.69923960063867)
```

Explanation: This data set was compiled from Julian Faraway’s data on head injury data. In order to get the data into a readable format for Stan (`.RDump`), Dr. Barber used the `rstan::stan_rdump` function within R.

For more information on creating RDump data files for use in Cmdstan, refer to: **RDump Format for CmdStan**.

Although JSON files were not used in this walkthrough, they should work similarly to RDump files, albeit the contents might be structured differently and the file extensions will be different when entering the data via the command line. For more information on creating JSON data files for use in Cmdstan, refer to: **JSON Format for CmdStan**.

Final steps

Now that all of the necessary files have been provided, the process remains relatively the same as before: `make`, `sample`, `display`.

To make the executable, open up a terminal window within the CmdStan directory (`~/cmdstan-2.29.2/`) and execute the following line:

```
make STANCLFLAGS=--allow-undefined USER_HEADER=examples/binomial_glm/external_binomial_glm.hpp \
  examples/binomial_glm/binomial_glm
```

If compilation was successful, you should find a file titled “`binomial_glm`” within your `/binomial_glm` directory.

Sample using the provided example data: `./binomial_glm sample data file=binGLM.Data.list.dump.R`

For a more interesting sampling routine, additional arguments can be provided to CmdStan and used within a for loop to produce multiple chains:

```
for i in {1..3}; do ./binomial_glm sample random seed=110710 \  
  data file=binGLM.Data.list.dump.R output refresh=1000 file=output_${i}.csv; done;
```

Display using the `stansummary` functionality: `~/cmdstan-2.29.2/bin/stansummary output_*.csv`

Additional points

Adding more inputs

Going off of the “Beyond the Bernoulli example” section above, you will likely need to add more than one input to the your external C++ code to get the most out of it. Assuming that the data type that you want to add is covered by Stan’s data types **here**, you should be able to adapt the previous code .stan and .cpp code chunks to include more inputs. In order to produce the templated code for these updated functions, you should follow the final bullet in the “External binomial hpp code” subsection.

Note: This is slightly more difficult in the case of matrices and vectors and these types are covered in the “Matrices and vertices” subsection below. Constrained data types were also not utilized in this walkthrough.

Including more files

Depending on the complexity of your external C++ code, it’s likely that you will need to include more than one C++ file within your CmdStan compilation. Luckily, the external C++ file that is supplied to CmdStan doesn’t have to be a header file (.hpp extension), but it can also be a source code file (.cpp). That means your external C++ file can follow this fairly common coding structure:

- Have a main.cpp file that is called during CmdStan compilation.
- Within that main.cpp file, you can include the headers for additional .hpp files.
- Within those .hpp files, you can include their associated .cpp files.

In this way, when CmdStan attempts to compile, the compiler will first look for main.cpp, then the .hpp files that main.cpp points to, and then the .cpp files that are pointed to by the .hpp files.

For example, instead of using the following **make** call:

```
make STANFLAGS=--allow-undefined USER_HEADER=examples/binomial_glm/external_binomial_glm.hpp \
  examples/binomial_glm/binomial_glm
```

... a user can use the following **make** call instead:

```
make STANFLAGS=--allow-undefined USER_HEADER=examples/binomial_glm/external_binomial_glm.cpp \
  examples/binomial_glm/binomial_glm
```

Within “external_binomial_glm.cpp” you could have an include statement at the top of the file that would link an additional header file. Like so:

```
#include "external_binomial_glm.hpp"
```

For more information on source file inclusion, refer to this **reference**.

Matrices and vertices

When it comes to matrix and vector multiplication, it is more difficult due to the fact that Stan uses the **Eigen library** to represent matrices and the **standard C++ library** to represent vectors. As such, your external C++ code will also need to use these libraries.

The difficulty comes with simple operations such as matrix multiplication, which are handled within the Eigen library. The Eigen functions would expect certain data types within the matrices. However, since those Eigen matrices contain Stan data types, you can’t use the functions from the Eigen library to handle the simple operations because compilation will fail. The C++ vectors behave in the same way.

The current work around to this issue is performing the matrix and vector operations within the Stan model code, and then storing the result into a vector within the Stan model code. The values can be iterated over afterwards, passed into an external function, and then the result is returned. This what was done for the binomial GLM example.

The Stan developers most likely have a solution to this issue. This most likely be in the form of Stan math functions that take the place of Eigen/C++ functions. Reaching out to them via the Stan discourse forum would be beneficial in figuring out this issue.

Limitations

The “limit” to CmdStan’s external C++ functionality is whether the custom function can be auto-differentiated by Stan. If not, then derivatives would need to be computed within the custom function.

However, aside from this unique limit, the typical Stan and C++ limitations still take exist. If your Stan model is not syntactically correct, then compilation will fail. If you C++ code is not syntactically correct, then compilation will fail.

It’s for these reasons that it might be easier to initially write code for the Stan model and C++ code separate from each other. Once both the Stan model and C++ code have been testing, the process for templating the C++ can be begin.

Acknowledgements

Dr. Jay Barber, the faculty member that oversaw this section of graduate research.

The Stan team for their **extensive documentation** and help via the **Stan discourse forums**.

Julian Faraway since the head injury data that used was sourced from their book, Linear Models with R.